

Homework 1: Gradient Descent & Friends

```
1 import pandas as pd
2 import numpy as np
3 import matplotlib.pyplot as plt
4
5 df = pd.read_csv("./real_estate.csv")
```

Question 1 (Pre-processing)

(a) Remove any rows of the data that contain a missing ('NA') value. List the indices of the removed data points. Then, delete all features from the dataset apart from: age, nearestMRT and nConvenience.

```
1 null_index = df[df.isnull().any(axis=1) == True].index
2 df.dropna(axis=0, inplace=True)
3 prices = df['price']
4 df.drop(columns=['transactiondate', 'latitude', 'longitude', 'price'],
5         inplace=True)
6 # print(null_index)
7 # print(df.head)
8 print("null indices: " , null_index.values)
```

```
1 null indices:  [ 19  41 109 144 230 301]
```

(b) normalisation and provide the mean value over your dataset

```
1 def pre_processing(x_data):
2     x_max = x_data.max()
3     x_min = x_data.min()
4     for _ in range(x_data.size):
5         _temp = (x_data[_] - x_min) / (x_max - x_min)
6         x_data[_] = _temp
7     return x_data
8
9 x_new_age = pre_processing(np.array(df['age']))
10 x_new_mrt = pre_processing(np.array(df['nearestMRT']))
```

```

11 x_new_nCon = pre_processing(np.array(df['nConvenience']))
12
13 x_new_age_mean = x_new_age.mean()
14 x_new_mrt_mean = x_new_mrt.mean()
15 x_new_nCon_mean = x_new_nCon.mean()
16
17 print("x_new_age_mean: ", x_new_age_mean)
18 print("x_new_nearestMRT_mean: ", x_new_mrt_mean)
19 print("x_new_nConvenience_mean: ", x_new_nCon_mean)

```

```

1 x_new_age_mean: 0.40607932670785213
2 x_new_nearestMRT_mean: 0.16264267697310722
3 x_new_nConvenience_mean: 0.4120098039215686

```

Question 2 (Train and Test sets)

first half of observations to create training set, remaining half for test set

```

1 x_new = pd.DataFrame(columns=['age', 'nearestMRT', 'nConvenience'])
2 x_new['age'] = x_new_age
3 x_new['nearestMRT'] = x_new_mrt
4 x_new['nConvenience'] = x_new_nCon
5 size = x_new.index.size
6 training_price = prices.values[:int(size / 2)]
7 test_price = prices.values[int(size / 2):]
8 training_set = x_new.values[:int(size / 2)]
9 test_set = x_new.values[int(size / 2):]

```

Print out the first and last rows of both your training and test sets

```

1 first_training_row = training_set[0]
2 last_training_row = training_set[-1]
3 first_test_row = test_set[0]
4 last_test_row = test_set[-1]
5
6 print("first training row: ", first_training_row)
7 print("last training row: ", last_training_row)
8 print("first test row: ", first_test_row)
9 print("last test row: ", last_test_row)

```

```
1 first training row: [0.73059361 0.00951267 1.          ]
2 last training row:  [0.87899543 0.09926012 0.3         ]
3 first test row:     [0.26255708 0.20677973 0.1         ]
4 last test row:      [0.14840183 0.0103754  0.9         ]
```

Question 3(Loss Function)

$$L_c(y, \hat{y}) = \frac{1}{n} \sum_{i=1}^n \left[\sqrt{\frac{1}{c^2} c y^{(i)} - \langle w^{(t)}, X^{(i)} \rangle^2 + 1} - 1 \right]$$

$$\frac{\partial L_c}{\partial w_k} = \frac{1}{n} \sum_{i=1}^n \left[\frac{1}{c^2} c y^{(i)} - \langle w^{(t)}, X^{(i)} \rangle^2 + 1 \right]^{\frac{1}{2}}$$

$$= \frac{\frac{\partial w^{(t)}}{\partial w_k} \cdot (-1) \cdot \frac{1}{c^2} \cdot c y^{(i)} - \langle w^{(t)}, X^{(i)} \rangle \cdot (2) \cdot \left(\frac{1}{2}\right)}{\sqrt{\frac{1}{c^2} c y^{(i)} - \langle w^{(t)}, X^{(i)} \rangle^2 + 1}}$$

$$= \frac{\frac{\partial w^{(t)}}{\partial w_k} \cdot \left(\langle w^{(t)}, X^{(i)} \rangle - y^{(i)} \right)}{c^2 \sqrt{\frac{1}{c^2} c y^{(i)} - \langle w^{(t)}, X^{(i)} \rangle^2 + 1}}$$

$$= \frac{\frac{\partial w^{(t)}}{\partial w_k} \cdot \left(\langle w^{(t)}, X^{(i)} \rangle - y^{(i)} \right)}{\sqrt{c^2 \left(\langle w^{(t)}, X^{(i)} \rangle - y^{(i)} \right)^2 + c^4}}$$

$$= \frac{\boxed{\frac{\partial w^{(t)}}{\partial w_k}} \cdot \left(\langle w^{(t)}, X^{(i)} \rangle - y^{(i)} \right)}{c^2 \sqrt{\frac{1}{c^2} \left(\langle w^{(t)}, X^{(i)} \rangle - y^{(i)} \right)^2 + 1}}$$

$$\text{for } \frac{\partial w^{(t)}}{\partial w_k} = \frac{\partial (w_0 + w_1 x_1^{(i)} + w_2 x_2^{(i)} + w_3 x_3^{(i)})}{\partial w_k}$$

$$= \frac{1 + x_1^{(i)} + x_2^{(i)} + x_3^{(i)}}{\downarrow}$$

let $= X_k^{(i)}$

$$\frac{x_k^{(i)} \cdot (\langle w^{(t)}, x^{(i)} \rangle - y^{(i)})}{c^2 \cdot \sqrt{\frac{1}{c^2} (\langle w^{(t)}, x^{(i)} \rangle - y^{(i)})^2 + 1}}$$

Question 4 (Gradient Descent Psuedocode)

gradient descent updates

```

1  for i <- 0 to nIteration:
2      derivate_sum <- 0
3      for j <- 0 to tranning_set_size:
4          derivate_sum += loss_function_derivate
5      derivate_loss_mean <- derivate_sum / tranning_set_size
6
7      w <- w - learning_rate * derivate_loss_mean

```

stochastic gradient descent updates

```
1 for i <- 0 to nIteration:
2     for epoch <- 0 to epoch_times:
3         w = w - learning_rate * loss_function_derivate
```

Question 5 (Gradient Descent Implementation)

helper function

```
1 def gradient_update(_X_k_i, w_t, training_values, prices_values, _size):
2     _sum_gd = 0
3     for _i in range(_size):
4         wX = np.dot(w_t, training_values[_i])
5         _sum_gd += ((_X_k_i * (wX - prices_values[_i])) / (2 * np.sqrt((wX
6 - prices_values[_i]) ** 2 + 1)))
7     return _sum_gd / _size
8
9 def loss_achieved(w_t, training_values, prices_values, _size, hyper):
10     _sum_loss = 0
11     for _i in range(_size):
12         wX = np.dot(w_t.T, training_values[_i])
13         _sum_loss += (np.sqrt(1 / (hyper ** 2) * (prices_values[_i] - wX)
14 ** 2 + 1) - 1)
15     return _sum_loss / _size
```

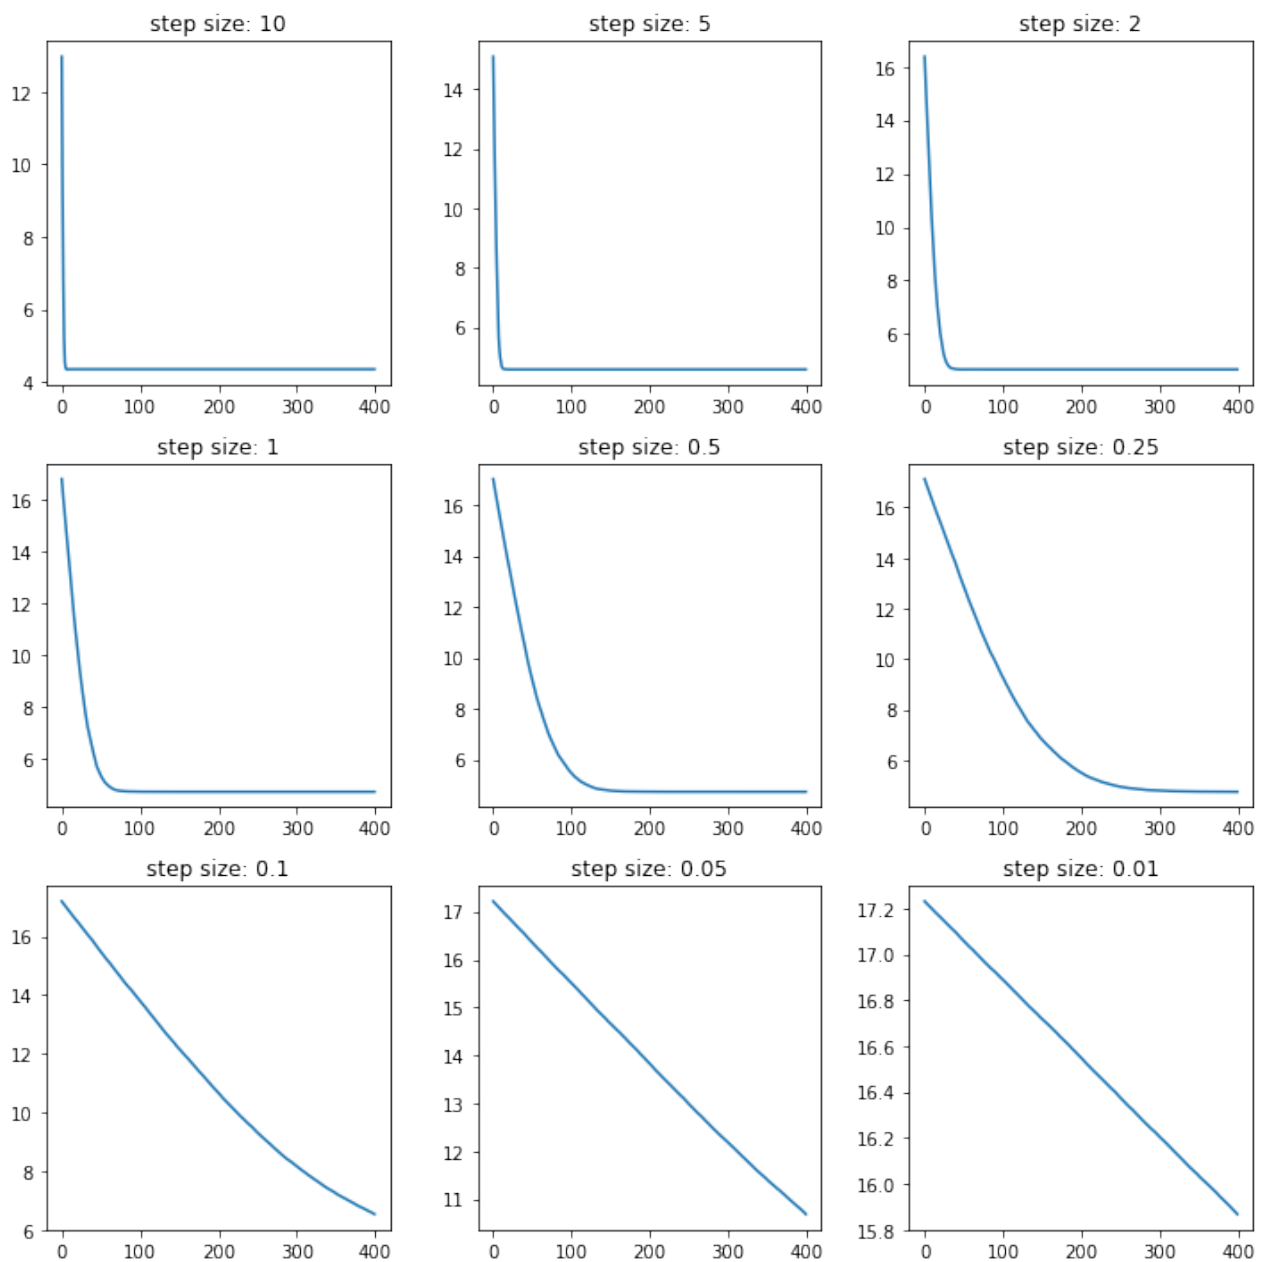
(a) Generate a 3×3 grid of plots showing performance for each step-size.

```
1 fig, ax = plt.subplots(3, 3, figsize=(10, 10))
2 nIter = 400
3 alphas = [10, 5, 2, 1, 0.5, 0.25, 0.1, 0.05, 0.01]
4 losses = []
5 training_size = training_set.shape[0]
6 training_gd_set = np.insert(training_set, 0, 1, axis=1)
7 c = 2
8 loss = []
9 w_plot = []
10 for i, ax in enumerate(ax.flat):
11     w = np.ones(4)
12     for index in range(nIter):
13         temp = index
14         if temp >= training_size:
15             temp -= training_size
```

```

16     w = w - gradient_update(training_gd_set[temp], w, training_gd_set,
17     training_price, training_size) * alphas[i]
17     loss_mean = loss_achieved(w, training_gd_set, training_price,
18     training_size, c)
18     loss.append(loss_mean)
19     losses.append(loss)
20     ax.plot(losses[i])
21     loss.clear()
22     ax.set_title(f"step size: {alphas[i]}")
23
24 plt.tight_layout()
25 plt.show()

```




(b) choose an appropriate step size (and state your choice), and explain why you made this choice.

I will choose **step size = 0.5**, for **0.1**, it is more stable and it does not fast to get the lowest loss.

(c) plot the progression of each of the four weights over the iterations, run your model on the train and test set, and print the achieved losses.

```
1  w_plot.clear()
2  fig, ax = plt.subplots(figsize=(10, 10))
3  w = np.ones(4)
4  w_plot.append(w)
5  for index in range(nIter):
6      temp = index
7      if temp >= training_size:
8          temp -= training_size
9      Xki = np.array([1, training_gd_set[temp][0], training_gd_set[temp][1],
10                     training_gd_set[temp][2]])
11     w = w - gradient_update(Xki, w, training_gd_set, training_price,
12                             training_size) * 0.3
13     w_plot.append(w)
14     ax.plot(w_plot)
15     plt.show()
16
17     print("w0: ", w[0])
18     print("w1: ", w[1])
19     print("w2: ", w[2])
20     print("w3: ", w[3])
21
22     test_size = test_set.shape[0]
23     test_set_gd_set = np.insert(test_set, 0, 1, axis=1)
24     test_loss = loss_achieved(w, test_set_gd_set, test_price, test_size, c)
25     train_loss = loss_achieved(w, training_gd_set, training_price,
26                                training_size, c)
27     print("batch training loss: ", train_loss)
28     print("batch test_loss: ", test_loss)
```

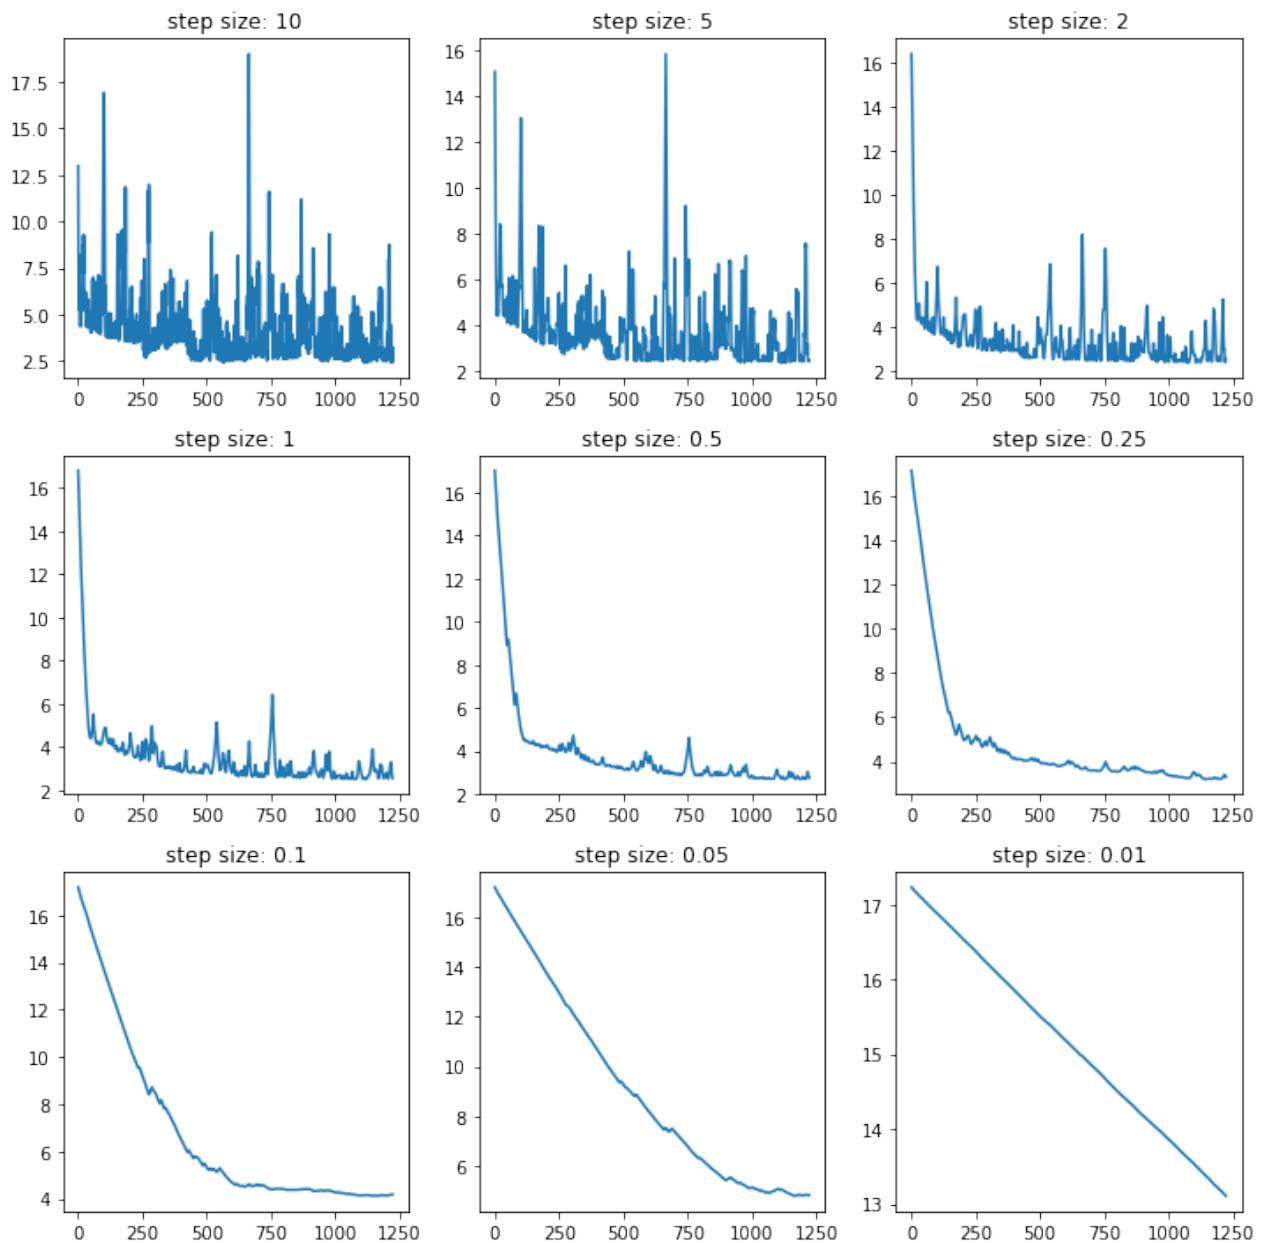



```
1 w0: 23.022722678965167
2 w1: 23.022722678965167
3 w2: 10.192369050047924
4 w3: 4.597963020259536
5 batch training loss: 6.097840919924557
6 batch test_loss: 5.516429121702068
```

Question 6

(a) plot Stochastic Gradient Descent Implementation

```
1  fig, ax = plt.subplots(3, 3, figsize=(10, 10))
2  losses.clear()
3  loss.clear()
4  epoch_times = 6
5  for i, ax in enumerate(ax.flat):
6      w = np.ones(4)
7      for index in range(training_size):
8          for _ in range(epoch_times):
9              # Xki = np.array([1, training_gd_set[index][0],
training_gd_set[index][1], training_gd_set[index][2]])
10             Xki = np.array(training_gd_set[index])
11             wX = np.dot(w, training_gd_set[index].T)
12             derivative_loss = (wX - training_price[index]) / (2 *
np.sqrt((wX - training_price[index]) ** 2 + 4))
13             w = w - alphas[i] * derivative_loss * Xki
14             loss_mean = loss_achieved(w, training_gd_set, training_price,
training_size, c)
15             loss.append(loss_mean)
16
17         losses.append(loss)
18         ax.plot(losses[i])
19         loss.clear()
20         ax.set_title(f"step size: {alphas[i]}")
21
22 plt.tight_layout()
23 plt.show()
```



(b) choose an appropriate step size, explain why

I will choose **step size = 0.25**, it is more stable and it can achieve lower loss.

(c) plot the progression of each of the four weights over the iterations, run your model on the train and test set, and record the achieved losses.

```


1 helper = []
2 w = np.ones(4)
3 helper.append(w)
4 for index in range(training_size):
5     for _ in range(epoch_times):
6         Xki = np.array(training_gd_set[index])

```

```

7         wX = np.dot(w, training_gd_set[index].T)
8         derivative_loss = (wX - training_price[index]) / (2 * np.sqrt((wX
- training_price[index]) ** 2 + 4))
9         w = w - 0.4 * derivative_loss * Xki
10        helper.append(w)
11
12    plt.plot(helper)
13    plt.show()
14
15    print("w0: ", w[0])
16    print("w1: ", w[1])
17    print("w2: ", w[2])
18    print("w3: ", w[3])
19
20    test_loss = loss_achieved(w, test_set_gd_set, test_price, test_size, c)
21    train_loss = loss_achieved(w, training_gd_set, training_price,
training_size, c)
22    print("SGD train_loss: ", train_loss)
23    print("SGD test_loss: ", test_loss)

```



```

1  w0:  29.209580076958723
2  w1:  -5.802603909995271
3  w2:  -11.093131590341907
4  w3:  24.654321874675617
5  SGD train_loss:  2.8777859585228605
6  SGD test_loss:   2.8711540128158446

```

Question7. Results Analysis

By adjusting different learning rate, we can get more stable and lower loss model. If we use a big learning rate, we may get a poor loss, however if we use a small learning rate, we need take more time to wait it to get the good loss.

Since GD need to get all sum and get the mean, it takes lots of time, however, sgd fix this problem.

SGD will frequently update gradient with a high variance causing fluctuation, therefore, GD has a smoother loss path than SGD.